

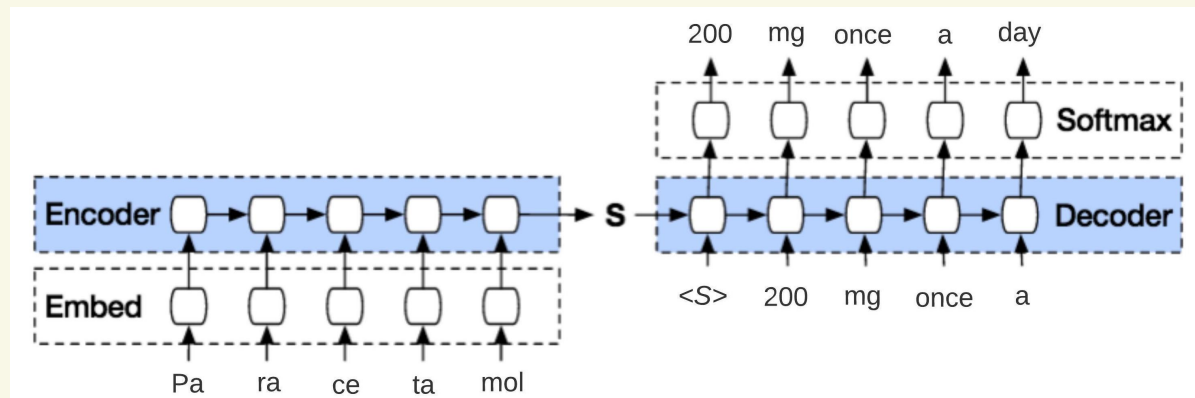
Generating Medical Prescriptions with Conditional Transformer

In: Workshop on Synthetic Data Generation with Generative AI at NeurIPS 2023, 16 Dec.

Motivations

- ▶ Access to real-world medication prescriptions is essential for medical research and healthcare quality improvement.
- ▶ This access is often limited due to the sensitive nature of the information expressed.
- ▶ Manually labelling these instructions for training and fine-tuning Natural Language Processing (NLP) models can be tedious and expensive.

LT3: Label to Text Transformer



LT3 Architecture with input/output behaviour (this is a shortened example of a generated synthetic medical prescription.)

Methodology

- ▶ A novel task-specific model architecture, **LT3: Label to Text Transformer**
- ▶ Tailored to generate synthetic medication prescriptions based on provided labels, e.g. a vocabulary list of medications and their attributes.
- ▶ LT3 is trained on a set of around 2K lines of medication prescriptions from the MIMIC-III
- ▶ Evaluate LT3's performance by contrasting it with T5, analysing the quality and diversity of generated texts.
- ▶ Deploy the generated synthetic data to train the SpacyNER model over the n2c2-2018 dataset.

Quality Evaluation to Reference and Intra-Diversity

Models	BLEU	ROUGE-1	ROUGE-2	ROUGE-L	BERTScore
T5 Small	71.75	76.16	66.24	75.55	0.70
T5 Base	71.98	76.28	66.30	75.45	0.70
T5 Large	69.89	75.07	65.19	74.22	0.68
LT3	78.52	78.16	68.72	77.55	0.72

Table: Closeness to Reference Evaluation with Lexical and Embedding Metrics

	Median Jaccard Score	Average Jaccard Score
LT3	0.650	0.652
T5 Base	0.658	0.660

Table: Jaccard scores of LT3 and T5 on Intra-diversity of generated text (lower is better)

Conventional BSD vs B2SD

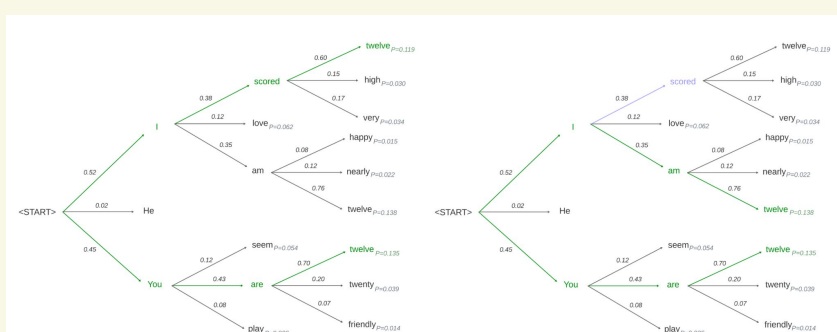



Figure 2: Execution Examples of Conventional Greedy BSD and B2SD Algorithms

- ▶ Beam Search Decoding using Backtracking (B2SD) to Transformers
- ▶ B2SD allows the model to find global optimised output instead of locally at each time step.

NER model training using real vs synthetic data



- ▶ Real data is over-sampled from 100 to 500%; RealSynthetic used 100% real plus synthetic data.
- ▶ The experiments show that the model trained on synthetic data can achieve a 96-98% F1 score at Label Recognition on Drug, Frequency, Route, Strength, and Form.
- ▶ **LT3** project including codes is **open-sourcelly hosted** at <https://github.com/HECTA-UoM/Label-To-Text-Transformer>

Samuel Belkadi [†], Nicolo Micheletti [†], Lifeng Han, Warren Del-Pinto, Goran Nenadic

{samuel.belkadi, nicolo.micheletti}@student.manchester.ac.uk | [†] co-first

Department of Computer Science, The University of Manchester, UK

